Marlos C. Machado
CMPUT 365: Introduction to Reinforcement Learning

Worksheet: Markov Decision Processes (MDPs)

Question 1. Assume you have a bandit problem with 4 actions, where the agent can see rewards from the set $\mathcal{R} = \{-3.0, -0.1, 0, 4.2\}$. Assume you have the probabilities for rewards for each action: $p(r|a)$ for $a \in \{1, 2, 3, 4\}$ and $r \in \{-3.0, -0.1, 0, 4.2\}$. How can you write this problem as an MDP? Remember that an MDP consists of $(\mathcal{S}, \mathcal{A}, \mathcal{R}, P, \gamma)$.

More abstractly, recall that a Bandit problem consists of a given action space $\mathcal{A} = \{1, ..., k\}$ (the k arms) and the distribution over rewards $p(r|a)$ for each action $a \in \mathcal{A}$. Specify an MDP that corresponds to this Bandit problem.

Question 2. Consider the 4×4 gridworld shown below.

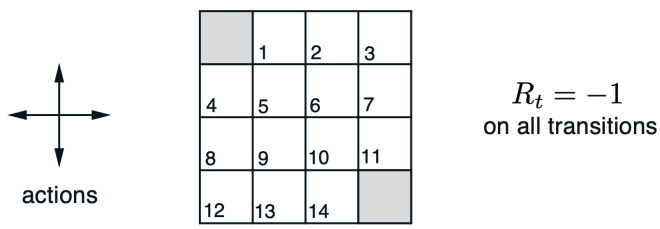

The nonterminal states are $\mathcal{S} = \{1, 2, \ldots, 14\}$. There are four actions possible in each state, $\mathcal{A} = \{$ `up`, `down`, `right`, `left` $\}$, which deterministically cause the corresponding state transitions, except that actions that would take the agent off the grid in fact leave the state unchanged. In this context, what are the values for the expressions below?

(a) $p(6, -1|5,\texttt{right}) =$ (b) $p(7, -1|7,\texttt{right}) =$ (c) $p(10, r|5,\texttt{right}) =$

Question 3. Prove that the discounted sum of rewards is always finite, if the rewards are bounded: $|R_{t+1}| \leq R_{\max}$ for all t for some finite $R_{\max} > 0$.

$$\left| \sum_{i=0}^{\infty} \gamma^i R_{t+1+i} \right| \qquad \text{for } \gamma \in [0, 1).$$

Hint: Recall that $|a + b| < |a| + |b|$.

Question 4. Exercise 3.6 from Sutton & Barto (2018).

Question 5. Exercise 3.7 from Sutton & Barto (2018).

Question 6. Exercise 3.8 from Sutton & Barto (2018).

Question 7. Exercise 3.9 from Sutton & Barto (2018).

Question 8. Exercise 3.10 from Sutton & Barto (2018).